

# Pet Image Recognition with Keras

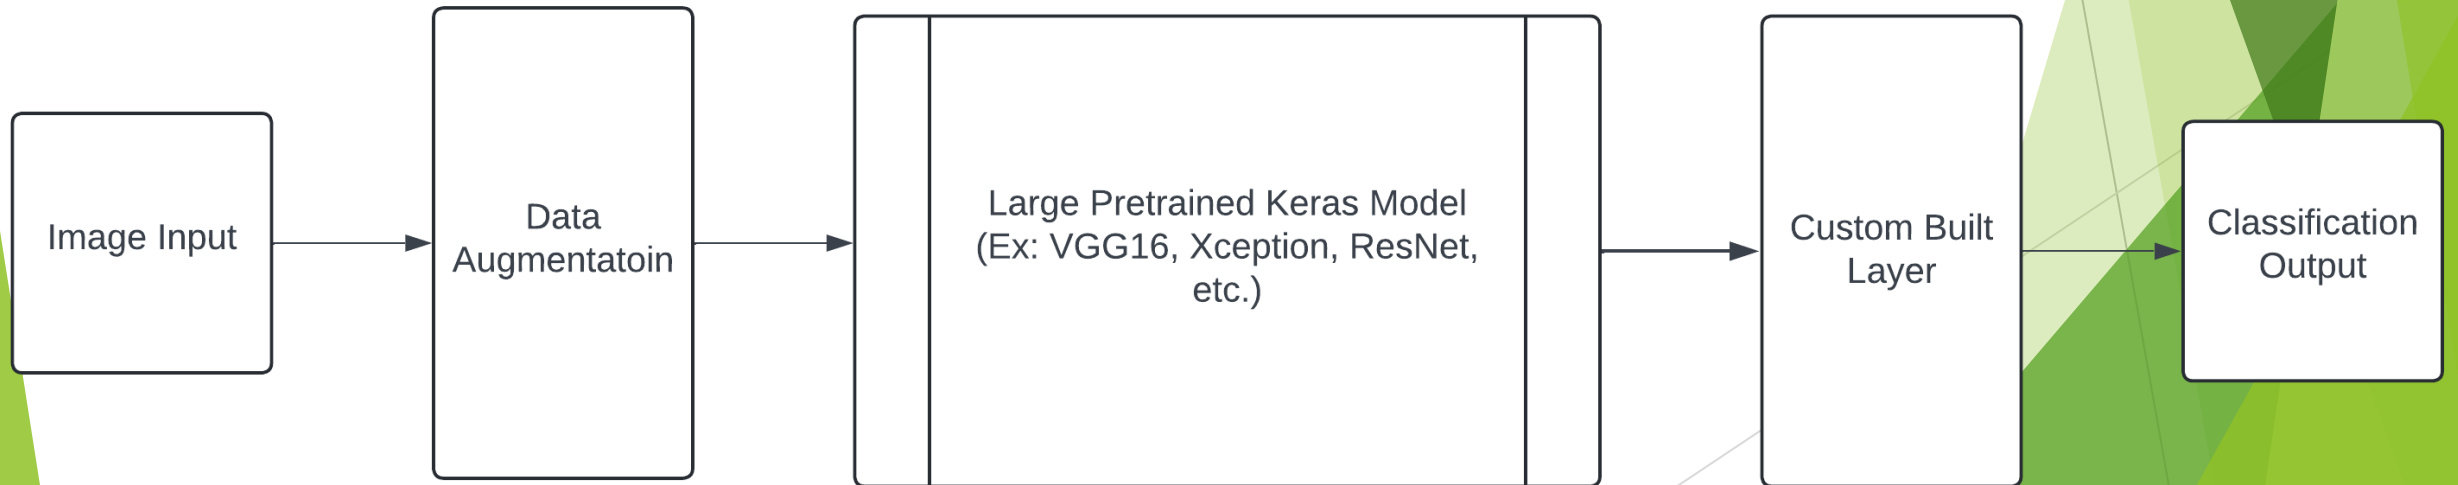
Evan Hollier and Paul Homuth

# Problem Statement

- ▶ Can a Image Recognition Model be used to classify 37 different pet breeds from photos?
  - ▶ What is the best pretrained model to use?
  - ▶ What are the best hyperparameters to add?


# Problem Setup

- Input: Image from Oxford IIIT Pet Dataset
- Early Layers: Large Pretrained Model from Keras
- Final Layers: Custom Built Layer tuned to our data
- Output: Classification of Pet




# Dataset Description

- ▶ Oxford IIIT Pet Dataset
  - ▶ A image set of 37 different breeds of pets. Each image has varying quality of light, scale, pose, and background making it a good random sampling of images
- ▶ Total of 7393 Images
  - ▶ 25 Dog breeds
  - ▶ 12 Cat breeds
  - ▶ Roughly 200 images per group



# Data Issues


- Damaged or corrupted images
  - Found a lot local to one class (Egyptian Mau)
  - Corrupted Images have missing/warped sections
- Confirming Total images
  - Reported 7349 images, but we found 7393 images
  - Caused a lot of indexing issues requiring cleanup



# Image Quality Overview

- ▶ Non standardize dataset
  - ▶ Varying backgrounds, light levels, distance to subject, and angle of photo
  - ▶ All ages of subject can be found. Puppies and Kittens included
- ▶ Below are examples from Abyssinian cat category

Light Variations



Background Variations



Age Variations




# Pretrained Model Selection

- ▶ Four selected pretrained models were tested to serve as a baseline
  - ▶ Xception, VGG\_16, ResNet, and EfficientNet
  - ▶ Selected due to wide range in complexity
- ▶ These baseline models were frozen and fit on our dataset
  - ▶ 37 Neurons, Softmax
- ▶ Baseline models were ran for 20 epochs
- ▶ Tracked the Loss and Accuracy metrics on both training and validation sets
  - ▶ Best model was selected out based on these metrics

# Pretrained Model Selection Cont.


VGG16

Training vs Validation Loss




ResNet50


Training vs Validation Loss



Training vs Validation Accuracy



Training vs Validation Accuracy







# Pretrained Model Selection Cont.

Xception

Training vs Validation Loss




Training vs Validation Accuracy




EfficientNetB5

Training vs Validation Loss



Training vs Validation Accuracy




# Grid Search

- ▶ Xception and EfficientNet baseline models needed a tie breaker
- ▶ Grid Searched across multiple parameters to determine winner in our custom layer
- ▶ Adjust Optimizer, Total number of Neurons, and Dropout Layer
  - ▶ Optimizers: RMSProp, Adam
  - ▶ Total Neurons: 16, 32, 64, 128
  - ▶ Dropout Layer: 0, .25, .5
- ▶ 10 epochs with early stopping patience of 4
- ▶ Best validation loss models were saved for comparison


# Best Xception vs Best EfficientNet

EfficientNet(n=32, d=0, opt=rmsprop)

Validation Loss

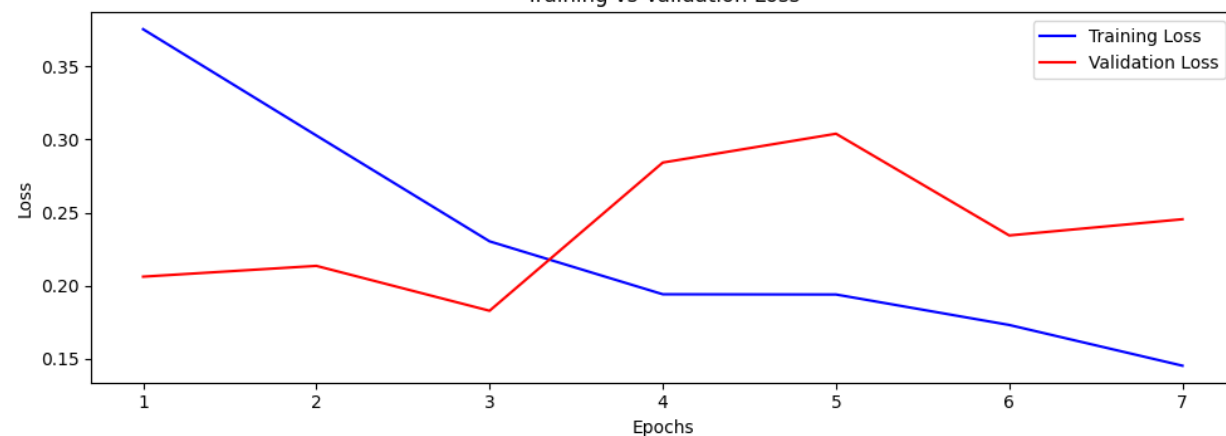


Validation Accuracy




Xcept(n=128, d=0.5, opt=adam)

Training vs Validation Loss




Training vs Validation Accuracy



# Further Custom Layer Tuning

- ▶ Found Xception performed best
  - ▶ Adam optimizer, 128 neurons, and .5 dropout
- ▶ Increased total training time
- ▶ Added more dense layers before output
  - ▶ Up to two additional layers
  - ▶ Tried more neurons (256, 384, 512)
  - ▶ .5 dropout
- ▶ Tried L1, L2, and L1\_L2 regularization
- ▶ Inserted Data Augmentation to increase training set
- ▶ Batch Normalization



# Final Model

- ▶ Input: Pet Image
- ▶ Data Augmentation: 1 out of 10,250 variations
- ▶ Early Layers: Xception with ImageNet weights
- ▶ Custom Layers: Two Densely Connected Layers
  - ▶ 256, 128 Neurons, Batch Normalized, 0.5 Dropouts
- ▶ Validation accuracy of 93.64%



# Results & Testing

- ▶ Test Accuracy: 90.3 % = 1337 / 1480
- ▶ Test Loss: 0.0174
- ▶ 37 Classes Tested with 40 samples each
  - ▶ 27 Found to have at least one mistake
  - ▶ Ragdoll and Russian Blue Cats had the most misclassifications
  - ▶ Dogs misclassified as cats: 1
  - ▶ Cats misclassified as dogs: 5
  - ▶ Three pairs that were confused 10+ times
- ▶ Potential Improvement calculated for each class

Class	Population %	Error %	Potential Improvement %
ragdoll	2.702703	30.000000	0.810811
russian_blue	2.702703	30.000000	0.810811
staffordshire_bull_terrier	2.702703	27.500000	0.743243
birman	2.702703	25.000000	0.675676
maine_coon	2.702703	25.000000	0.675676

ragdoll	birman	x16
american_pit_bull_terrier	staffordshire_bull_terrier	x11
egyptian_mau	bengal	x11
british_shorthair	russian_blue	x8
russian_blue	bombay	x8
siamese	birman	x7
american_pit_bull_terrier	american_bulldog	x4
bengal	maine_coon	x4

True Label: ragdoll	
birman	x9
persian	x1
samoyed	x1
siamese	x1

True Label: russian_blue	
bombay	x8
british_shorthair	x2
abyssinian	x1
bengal	x1

# Most Confused Classes

True: ragdoll




Predicted: birman



True: american\_pit\_bull\_terrier



Predicted: staffordshire\_bull\_terrier



True: egyptian\_mau



Predicted: bengal





# Testing on Our Own Pets

Theo: Maine coon / Domestic Long Hair mix

Pred: maine\_coon

Pred: maine\_coon

Pred: maine\_coon



Pepper: Russian blue / Tortoiseshell mix

Pred: russian\_blue



Pred: russian\_blue


Dude: Wheaten terrier

Pred: wheaten\_terrier

Pred: wheaten\_terrier



Pred: abyssinian





# Conclusion

- ▶ A image recognition model was trained and built:
  - ▶ Used Data Augmentation during training
  - ▶ Xception with ImageNet weights
  - ▶ Two Densely Connected Layers (256, 128)
  - ▶ Batch Normalization
  - ▶ Dropouts and variable learning rate
- ▶ The model produced a validation accuracy of 93.64 and a testing accuracy of 90.34
- ▶ Using these techniques is feasible to build an image recognition model to distinguish different pets